

# Analyse eines Forschungsthemas

## Stochastic Shortest Paths

Maximilian Starke

Fakultät für Informatik  
Technische Universität Dresden

3. Januar 2023

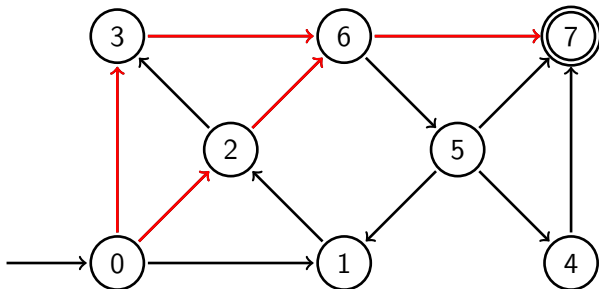
# Section 1

## Introduction






- The *simplest* shortest path problem

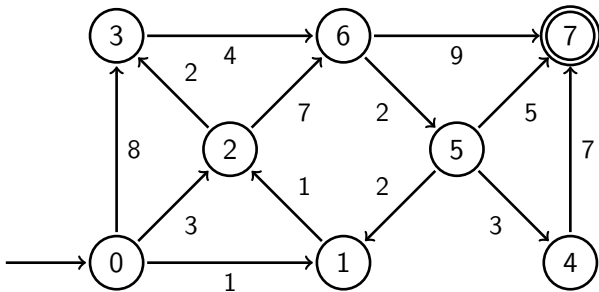


- Task
  - Find the shortest path (*number of hops*)!

- The *classical, non-stochastic, deterministic* shortest path problem

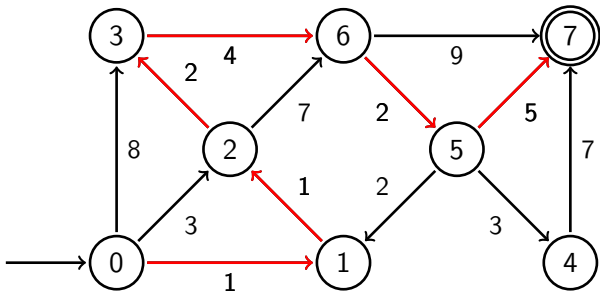


- The *classical, non-stochastic, deterministic* shortest path problem



- Task
  - Find the path with the minimal weight sum!

- The *classical, non-stochastic, deterministic* shortest path problem

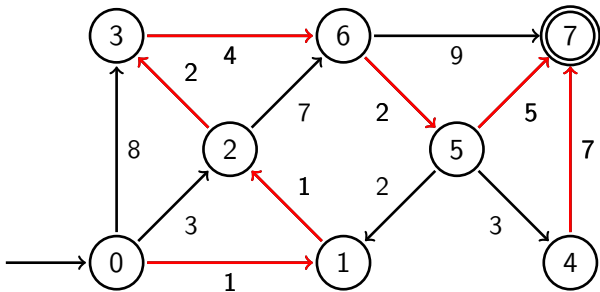


- Task
  - Find the path with the minimal weight sum!



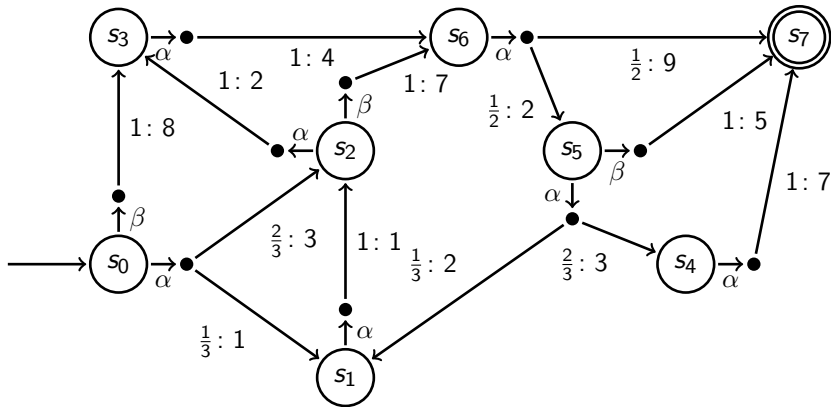


- The *classical, non-stochastic, deterministic* shortest path problem

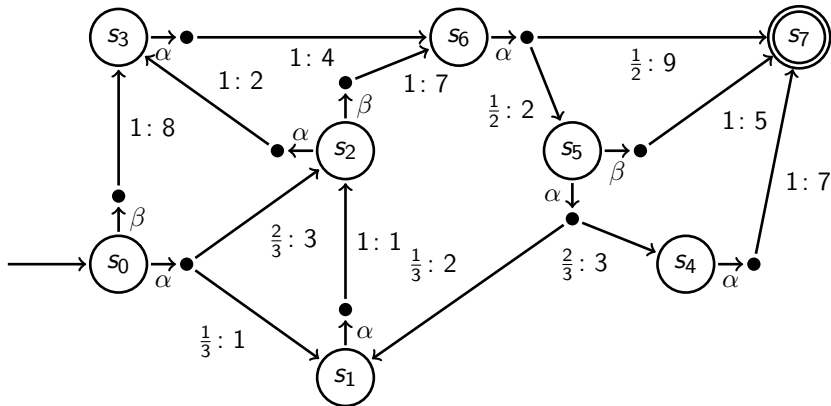


- Task
  - Find the path with the minimal weight sum!
  - Give a strategy to always reach the goal while collecting minimal weight!

► The *stochastic* shortest path problem

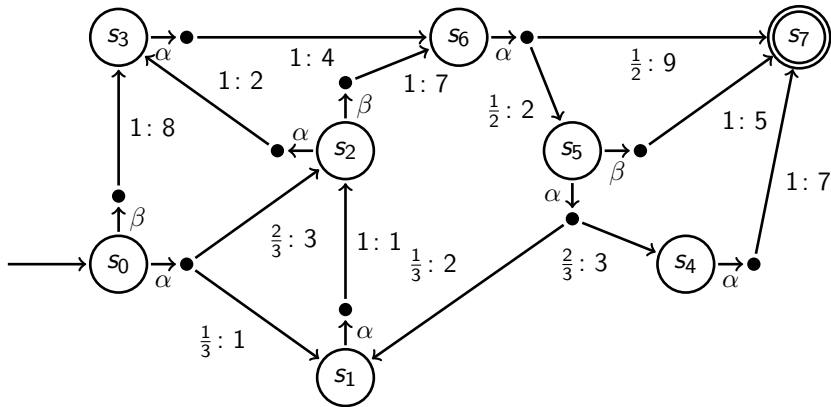


► The *stochastic* shortest path problem



► Markov Decision Process (MDP)

► The *stochastic* shortest path problem



► Markov Decision Process (MDP)

► Task

- Give a strategy to reach the goal with minimal *expected* accumulated weights!

# MDP - Markov Decision Process (I)

## definition

An MDP is a tuple

$$\mathcal{M} = (S, Act, P, s_{init}, wgt)$$

# MDP - Markov Decision Process (I)

## definition

An MDP is a tuple

$$\mathcal{M} = (S, Act, P, s_{init}, wgt)$$

where

- ▶  $S$  is a finite set of states.

# MDP - Markov Decision Process (I)

## definition

An MDP is a tuple

$$\mathcal{M} = (S, Act, P, s_{init}, wgt)$$

where

- ▶  $S$  is a finite set of states.
- ▶  $Act$  is a finite set of actions.



# MDP - Markov Decision Process (I)

## definition

An MDP is a tuple

$$\mathcal{M} = (S, Act, P, s_{init}, wgt)$$

where

- ▶  $S$  is a finite set of states.
- ▶  $Act$  is a finite set of actions.
- ▶  $P : S \times Act \dashrightarrow Distr(S)$  maps every enabled state - action - pair  $(s, \alpha)$  to a probability distribution over the states.

# MDP - Markov Decision Process (I)

## definition

An MDP is a tuple

$$\mathcal{M} = (S, Act, P, s_{init}, wgt)$$

where

- ▶  $S$  is a finite set of states.
- ▶  $Act$  is a finite set of actions.
- ▶  $P : S \times Act \dashrightarrow Distr(S)$  maps every enabled state - action - pair  $(s, \alpha)$  to a probability distribution over the states.
- ▶  $s_{init} \in S$  is some designated initial state.

# MDP - Markov Decision Process (I)

## definition

An MDP is a tuple

$$\mathcal{M} = (S, Act, P, s_{init}, wgt)$$

where

- ▶  $S$  is a finite set of states.
- ▶  $Act$  is a finite set of actions.
- ▶  $P : S \times Act \dashrightarrow Distr(S)$  maps every enabled state - action - pair  $(s, \alpha)$  to a probability distribution over the states.
- ▶  $s_{init} \in S$  is some designated initial state.
- ▶  $wgt : S \times Act \rightarrow \mathbb{Z}$  is some weight (reward) function.

# MDP - Markov Decision Process (II)

## Paths

$$\pi \in \text{Paths}(\mathcal{M})$$

# MDP - Markov Decision Process (II)

## Paths

$$\pi \in \text{Paths}(\mathcal{M})$$

►  $\pi = (s_0, \alpha_0, s_1, \alpha_1, \dots)$

# MDP - Markov Decision Process (II)

## Paths

$$\pi \in \text{Paths}(\mathcal{M})$$

- ▶  $\pi = (s_0, \alpha_0, s_1, \alpha_1, \dots)$
- ▶  $\pi$  is some maximal finite or infinite path

# MDP - Markov Decision Process (II)

## Paths

$$\pi \in \text{Paths}(\mathcal{M})$$

- ▶  $\pi = (s_0, \alpha_0, s_1, \alpha_1, \dots)$
- ▶  $\pi$  is some maximal finite or infinite path

Let  $F \subseteq S$  be some set of goal states...

# MDP - Markov Decision Process (II)

## Paths

$$\pi \in \text{Paths}(\mathcal{M})$$

- ▶  $\pi = (s_0, \alpha_0, s_1, \alpha_1, \dots)$
- ▶  $\pi$  is some maximal finite or infinite path

Let  $F \subseteq S$  be some set of goal states...

accumulated weights until reaching a goal

$$\Diamond F : \text{Paths}(\mathcal{M}^{\mathfrak{S}}) \rightarrow \mathbb{Q} :$$

$$\Diamond F(\pi) = \begin{cases} \text{wgt}(\hat{\pi}), & \hat{\pi} \text{ is shortest prefix of } \pi \text{ s.t. } \text{last}(\hat{\pi}) \in F \\ \infty, & \text{otherwise, i.e. } \pi \not\models \Diamond F \end{cases}$$



## Section 2

### The classic stochastic shortest path problem

# The classic stochastic shortest path problem

- ▶ given:
  - ▶ a single goal state

# The classic stochastic shortest path problem

- ▶ given:
  - ▶ a single goal state
  - ▶ positive cycle condition: There is no cycle  $\pi$  with  $\text{wgt}(\pi) \leq 0$

# The classic stochastic shortest path problem

- ▶ given:
  - ▶ a single goal state
  - ▶ positive cycle condition: There is no cycle  $\pi$  with  $\text{wgt}(\pi) \leq 0$
  - ▶ goal is reachable from each state

# The classic stochastic shortest path problem

- ▶ given:
  - ▶ a single goal state
  - ▶ positive cycle condition: There is no cycle  $\pi$  with  $\text{wgt}(\pi) \leq 0$
  - ▶ goal is reachable from each state
- ▶ objective: Minimize the expected accumulated weight until reaching goal state.

$$\mathbb{E}_{\mathcal{M},s}^{\text{inf}}(\Diamond \text{goal}) := \inf_{\mathfrak{S} \in \text{Schedulers}(M)} \mathbb{E}_{\mathcal{M},s}^{\mathfrak{S}}(\Diamond \text{goal})$$

# The classic stochastic shortest path problem

- ▶ given:
  - ▶ a single goal state
  - ▶ positive cycle condition: There is no cycle  $\pi$  with  $\text{wgt}(\pi) \leq 0$
  - ▶ goal is reachable from each state
- ▶ objective: Minimize the expected accumulated weight until reaching goal state.

$$\mathbb{E}_{\mathcal{M},s}^{\text{inf}}(\Diamond \text{goal}) := \inf_{\mathfrak{S} \in \text{Schedulers}(M)} \mathbb{E}_{\mathcal{M},s}^{\mathfrak{S}}(\Diamond \text{goal})$$

- ▶ Well known for a long time:
  - ▶ There exists an optimal memoryless deterministic scheduler  $\mathfrak{S}$ .

# The classic stochastic shortest path problem

- ▶ given:
  - ▶ a single goal state
  - ▶ positive cycle condition: There is no cycle  $\pi$  with  $\text{wgt}(\pi) \leq 0$
  - ▶ goal is reachable from each state
- ▶ objective: Minimize the expected accumulated weight until reaching goal state.

$$\mathbb{E}_{\mathcal{M},s}^{\text{inf}}(\Diamond \text{goal}) := \inf_{\mathfrak{S} \in \text{Schedulers}(M)} \mathbb{E}_{\mathcal{M},s}^{\mathfrak{S}}(\Diamond \text{goal})$$

- ▶ Well known for a long time:
  - ▶ There exists an optimal memoryless deterministic scheduler  $\mathfrak{S}$ .
  - ▶  $\mathfrak{S}$  is computable by solving a LP

# The classic stochastic shortest path problem

- ▶ given:
  - ▶ a single goal state
  - ▶ positive cycle condition: There is no cycle  $\pi$  with  $\text{wgt}(\pi) \leq 0$
  - ▶ goal is reachable from each state
- ▶ objective: Minimize the expected accumulated weight until reaching goal state.

$$\mathbb{E}_{\mathcal{M},s}^{\text{inf}}(\Diamond \text{goal}) := \inf_{\mathfrak{S} \in \text{Schedulers}(M)} \mathbb{E}_{\mathcal{M},s}^{\mathfrak{S}}(\Diamond \text{goal})$$

- ▶ Well known for a long time:
  - ▶ There exists an optimal memoryless deterministic scheduler  $\mathfrak{S}$ .
  - ▶  $\mathfrak{S}$  is computable by solving a LP
  - ▶ iterative algorithm:



# The classic stochastic shortest path problem

- ▶ given:
  - ▶ a single goal state
  - ▶ positive cycle condition: There is no cycle  $\pi$  with  $\text{wgt}(\pi) \leq 0$
  - ▶ goal is reachable from each state
- ▶ objective: Minimize the expected accumulated weight until reaching goal state.

$$\mathbb{E}_{\mathcal{M},s}^{\text{inf}}(\Diamond \text{goal}) := \inf_{\mathfrak{S} \in \text{Schedulers}(M)} \mathbb{E}_{\mathcal{M},s}^{\mathfrak{S}}(\Diamond \text{goal})$$

- ▶ Well known for a long time:
  - ▶ There exists an optimal memoryless deterministic scheduler  $\mathfrak{S}$ .
  - ▶  $\mathfrak{S}$  is computable by solving a LP
  - ▶ iterative algorithm:
    - ▶ start at any feasible scheduler

# The classic stochastic shortest path problem

- ▶ given:
  - ▶ a single goal state
  - ▶ positive cycle condition: There is no cycle  $\pi$  with  $\text{wgt}(\pi) \leq 0$
  - ▶ goal is reachable from each state
- ▶ objective: Minimize the expected accumulated weight until reaching goal state.

$$\mathbb{E}_{\mathcal{M},s}^{\text{inf}}(\Diamond \text{goal}) := \inf_{\mathfrak{S} \in \text{Schedulers}(M)} \mathbb{E}_{\mathcal{M},s}^{\mathfrak{S}}(\Diamond \text{goal})$$

- ▶ Well known for a long time:
  - ▶ There exists an optimal memoryless deterministic scheduler  $\mathfrak{S}$ .
  - ▶  $\mathfrak{S}$  is computable by solving a LP
  - ▶ iterative algorithm:
    - ▶ start at any feasible scheduler
    - ▶ iterative improvement

# The classic stochastic shortest path problem

- ▶ given:
  - ▶ a single goal state
  - ▶ positive cycle condition: There is no cycle  $\pi$  with  $\text{wgt}(\pi) \leq 0$
  - ▶ goal is reachable from each state
- ▶ objective: Minimize the expected accumulated weight until reaching goal state.

$$\mathbb{E}_{\mathcal{M},s}^{\text{inf}}(\Diamond \text{goal}) := \inf_{\mathfrak{S} \in \text{Schedulers}(M)} \mathbb{E}_{\mathcal{M},s}^{\mathfrak{S}}(\Diamond \text{goal})$$

- ▶ Well known for a long time:
  - ▶ There exists an optimal memoryless deterministic scheduler  $\mathfrak{S}$ .
  - ▶  $\mathfrak{S}$  is computable by solving a LP
  - ▶ iterative algorithm:
    - ▶ start at any feasible scheduler
    - ▶ iterative improvement
    - ▶ stop at an optimal vertex of the LP (corresponding to some MD scheduler)

# The classic stochastic shortest path problem

- ▶ Can we do it better?

# The classic stochastic shortest path problem

- ▶ Can we do it better?  $\longrightarrow$  YES! - using **spider construction**!

# The classic stochastic shortest path problem

- ▶ Can we do it better?  $\longrightarrow$  **YES! - using spider construction!**
- ▶ given:
  - ▶ a single goal state
  - ▶ positive cycle condition: There is no cycle  $\pi$  with  $\text{wgt}(\pi) \leq 0$

# The classic stochastic shortest path problem

- ▶ Can we do it better? → **YES! - using spider construction!**
- ▶ given:
  - ▶ a single goal state
  - ▶  $\mathcal{M}$  is an MDP with arbitrary integer weights

# The classic stochastic shortest path problem

- ▶ Can we do it better? → **YES! - using spider construction!**
- ▶ given:
  - ▶ a single goal state
  - ▶  $\mathcal{M}$  is an MDP with arbitrary integer weights
  - ▶ goal is reachable from each state



# The classic stochastic shortest path problem

- ▶ Can we do it better?  $\longrightarrow$  **YES! - using spider construction!**
- ▶ given:
  - ▶ a single goal state
  - ▶  $\mathcal{M}$  is an MDP with arbitrary integer weights
  - ▶ goal is reachable from each state
- ▶ The following can be solved in polynomial time:
  - ▶ Check:  $\mathbb{E}_{\mathcal{M},s}^{inf}(\Diamond \text{goal}) > -\infty$ ?
  - ▶ Compute  $\mathbb{E}_{\mathcal{M},s}^{inf}$  if it is finite

# Spider Construction

- ▶ Idea: construct a new MDP  $\mathcal{N}$  from the given MDP  $\mathcal{M}$

# Spider Construction

- ▶ Idea: construct a new MDP  $\mathcal{N}$  from the given MDP  $\mathcal{M}$
- ▶ Pick a 0-BSCC  $\mathcal{E}$  of  $\mathcal{M}$  and some vertex  $s_0$  in  $\mathcal{E}$ .

# Spider Construction

- ▶ Idea: construct a new MDP  $\mathcal{N}$  from the given MDP  $\mathcal{M}$
- ▶ Pick a 0-BSCC  $\mathcal{E}$  of  $\mathcal{M}$  and some vertex  $s_0$  in  $\mathcal{E}$ .
- ▶  $\mathcal{M} \mapsto \mathcal{N} := \text{Spider}_{\mathcal{E}, s_0}(\mathcal{M})$

# Spider Construction

- ▶ Idea: construct a new MDP  $\mathcal{N}$  from the given MDP  $\mathcal{M}$
- ▶ Pick a 0-BSCC  $\mathcal{E}$  of  $\mathcal{M}$  and some vertex  $s_0$  in  $\mathcal{E}$ .
- ▶  $\mathcal{M} \mapsto \mathcal{N} := \text{Spider}_{\mathcal{E}, s_0}(\mathcal{M})$
- ▶ The spider construction is done by applying the following steps:
  1. Remove all actions  $(s, \alpha_s) \in \mathcal{E}$

# Spider Construction

- ▶ Idea: construct a new MDP  $\mathcal{N}$  from the given MDP  $\mathcal{M}$
- ▶ Pick a 0-BSCC  $\mathcal{E}$  of  $\mathcal{M}$  and some vertex  $s_0$  in  $\mathcal{E}$ .
- ▶  $\mathcal{M} \mapsto \mathcal{N} := \text{Spider}_{\mathcal{E}, s_0}(\mathcal{M})$
- ▶ The spider construction is done by applying the following steps:
  1. Remove all actions  $(s, \alpha_s) \in \mathcal{E}$
  2. Add actions  $(s, \tau)$  for all  $s \in \mathcal{E} \setminus \{s_0\}$  such that
    - ▶  $P_{\mathcal{N}}(s, \tau, s_0) := 1$

# Spider Construction

- ▶ Idea: construct a new MDP  $\mathcal{N}$  from the given MDP  $\mathcal{M}$
- ▶ Pick a 0-BSCC  $\mathcal{E}$  of  $\mathcal{M}$  and some vertex  $s_0$  in  $\mathcal{E}$ .
- ▶  $\mathcal{M} \mapsto \mathcal{N} := \text{Spider}_{\mathcal{E}, s_0}(\mathcal{M})$
- ▶ The spider construction is done by applying the following steps:
  1. Remove all actions  $(s, \alpha_s) \in \mathcal{E}$
  2. Add actions  $(s, \tau)$  for all  $s \in \mathcal{E} \setminus \{s_0\}$  such that
    - ▶  $P_{\mathcal{N}}(s, \tau, s_0) := 1$
    - ▶  $\text{wgt}_{\mathcal{N}}(s, \tau) := \text{wgt}(s, s_0)$

# Spider Construction

- ▶ Idea: construct a new MDP  $\mathcal{N}$  from the given MDP  $\mathcal{M}$
- ▶ Pick a 0-BSCC  $\mathcal{E}$  of  $\mathcal{M}$  and some vertex  $s_0$  in  $\mathcal{E}$ .
- ▶  $\mathcal{M} \mapsto \mathcal{N} := \text{Spider}_{\mathcal{E}, s_0}(\mathcal{M})$
- ▶ The spider construction is done by applying the following steps:
  1. Remove all actions  $(s, \alpha_s) \in \mathcal{E}$
  2. Add actions  $(s, \tau)$  for all  $s \in \mathcal{E} \setminus \{s_0\}$  such that
    - ▶  $P_{\mathcal{N}}(s, \tau, s_0) := 1$
    - ▶  $\text{wgt}_{\mathcal{N}}(s, \tau) := \text{wgt}(s, s_0)$
  3. For each  $s \in \mathcal{E} \setminus \{s_0\}$  and  $\beta \in \text{Act}_{\mathcal{M}}(s) \setminus \{\alpha_s\}$  let us replace  $(s, \beta)$  by  $(s_0, \beta)$  where
    - ▶  $P_{\mathcal{N}}(s_0, \beta, u) := P_{\mathcal{M}}(s, \beta, u)$
    - ▶  $\text{wgt}_{\mathcal{N}}(s_0, \beta) + \text{wgt}(s, s_0) = \text{wgt}_{\mathcal{M}}(s, \beta)$



# Classification of paths

A path  $\pi \in \text{InfPaths}(\mathcal{M})$  is called

- ▶ pumping  $:\Leftrightarrow \liminf_{n \rightarrow \infty} (\text{wgt}(\text{pref}(\pi, n))) = \infty$

# Classification of paths

A path  $\pi \in \text{InfPaths}(\mathcal{M})$  is called

- ▶ pumping  $:\Leftrightarrow \liminf_{n \rightarrow \infty} (\text{wgt}(\text{pref}(\pi, n))) = \infty$
- ▶ (positively)  
negatively weight divergent  $:\Leftrightarrow \begin{array}{l} \limsup_{n \rightarrow \infty} \\ \liminf_{n \rightarrow \infty} \end{array} = \begin{array}{l} \infty \\ -\infty \end{array}$

# Classification of paths

A path  $\pi \in \text{InfPaths}(\mathcal{M})$  is called

- ▶ pumping  $:\Leftrightarrow \liminf_{n \rightarrow \infty} (\text{wgt}(\text{pref}(\pi, n))) = \infty$
- ▶ (positively)  
negatively weight divergent  $:\Leftrightarrow \begin{array}{l} \limsup_{n \rightarrow \infty} \\ \liminf_{n \rightarrow \infty} \end{array} = \begin{array}{l} \infty \\ -\infty \end{array}$
- ▶ gambling  $:\Leftrightarrow \pi$  is positively and negatively weight divergent

# Classification of paths

A path  $\pi \in \text{InfPaths}(\mathcal{M})$  is called

- ▶ pumping  $:\Leftrightarrow \liminf_{n \rightarrow \infty} (\text{wgt}(\text{pref}(\pi, n))) = \infty$
- ▶ (positively)  
negatively weight divergent  $:\Leftrightarrow \begin{array}{l} \limsup_{n \rightarrow \infty} \\ \liminf_{n \rightarrow \infty} \end{array} = \begin{array}{l} \infty \\ -\infty \end{array}$
- ▶ gambling  $:\Leftrightarrow \pi$  is positively and negatively weight divergent
- ▶ bounded from below  $:\Leftrightarrow \liminf_{n \rightarrow \infty} \text{wgt}(\text{pref}(\pi, n)) \in \mathbb{Z}$

# Classification of end components

We distinguish end components by the following types

- ▶ pumping ECs:  $\exists$  scheduler  $\mathcal{G} : \Pr(\pi \text{ is pumping}) = 1$

# Classification of end components

We distinguish end components by the following types

► pumping ECs:  $\exists$  scheduler  $\mathfrak{S} : \Pr(\pi \text{ is pumping}) = 1$

► (positively)  
negatively weight divergent ECs:

$\exists$  scheduler  $\mathfrak{S} : \Pr\left(\pi \text{ is } \begin{matrix} \text{(positively)} \\ \text{negatively} \end{matrix} \text{ weight divergent} \right) = 1$

# Classification of end components

We distinguish end components by the following types

- ▶ pumping ECs:  $\exists$  scheduler  $\mathfrak{S} : \Pr(\pi \text{ is pumping}) = 1$
- ▶ (positively)  
negatively weight divergent ECs:  
 $\exists$  scheduler  $\mathfrak{S} : \Pr(\pi \text{ is } \begin{smallmatrix} \text{(positively)} \\ \text{negatively} \end{smallmatrix} \text{ weight divergent}) = 1$
- ▶ gambling ECs: There is a scheduler s.t.  $\mathbb{E}(\text{MP}) = 0$  and it is positively and negatively weight divergent

# Classification of end components

We distinguish end components by the following types

- ▶ pumping ECs:  $\exists$  scheduler  $\mathfrak{S} : \Pr(\pi \text{ is pumping}) = 1$
- ▶ (positively)  
negatively weight divergent ECs:  
 $\exists$  scheduler  $\mathfrak{S} : \Pr(\pi \text{ is } \begin{smallmatrix} \text{(positively)} \\ \text{negatively} \end{smallmatrix} \text{ weight divergent}) = 1$
- ▶ gambling ECs: There is a scheduler s.t.  $\mathbb{E}(\text{MP}) = 0$  and it is positively and negatively weight divergent
- ▶ bounded EC: There exists an upper bound and a lower bound



## Check Weight-Divergence of a SCC

- ▶ can be done in **PTime**

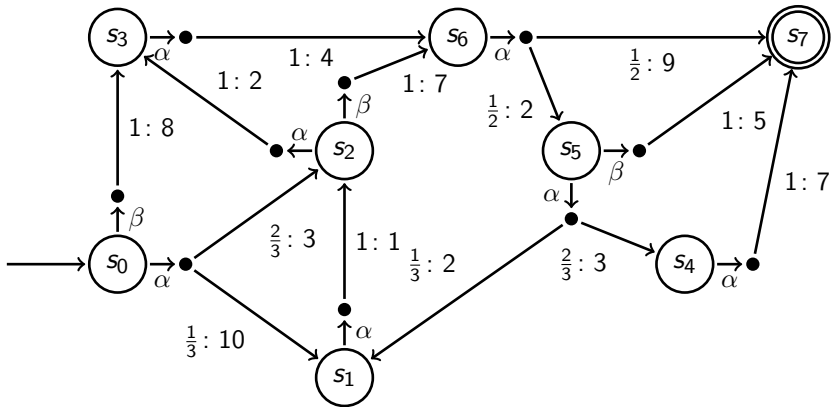
## Check Weight-Divergence of a SCC

- ▶ can be done in **PTime**
- ▶ two cases:
  - ▶ "yes, weight-divergent": finds a gambling or a pumping scheduler

## Check Weight-Divergence of a SCC

- ▶ can be done in **PTime**
- ▶ two cases:
  - ▶ "yes, weight-divergent": finds a gambling or a pumping scheduler
  - ▶ "no": returns an *equivalent* MDP  $\mathcal{N}$  without 0-ECs

- ▶ The *stochastic* shortest path problem

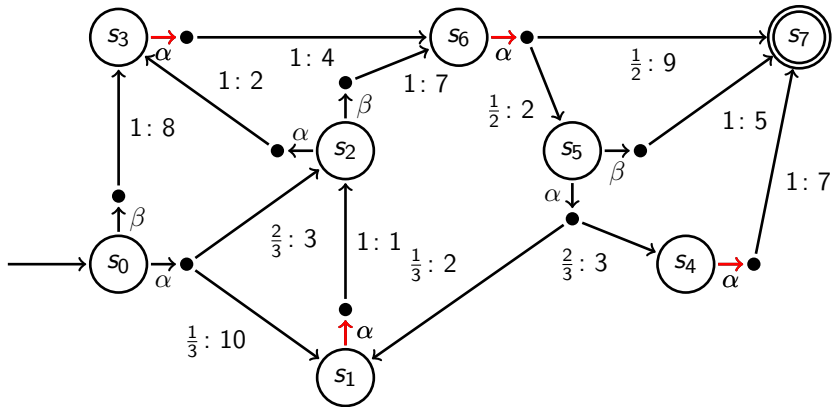


- ▶ Markov Decision Process (MDP)

► Task

- ▶ Give a strategy to reach the goal with minimal *expected* accumulated weights!

- ▶ The *stochastic* shortest path problem



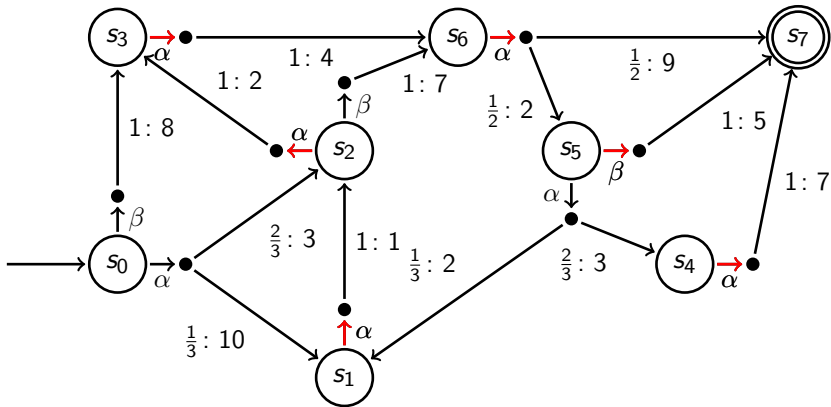
- ▶ Markov Decision Process (MDP)

► Task

- ▶ Give a strategy to reach the goal with minimal *expected* accumulated weights!



- ▶ The *stochastic* shortest path problem

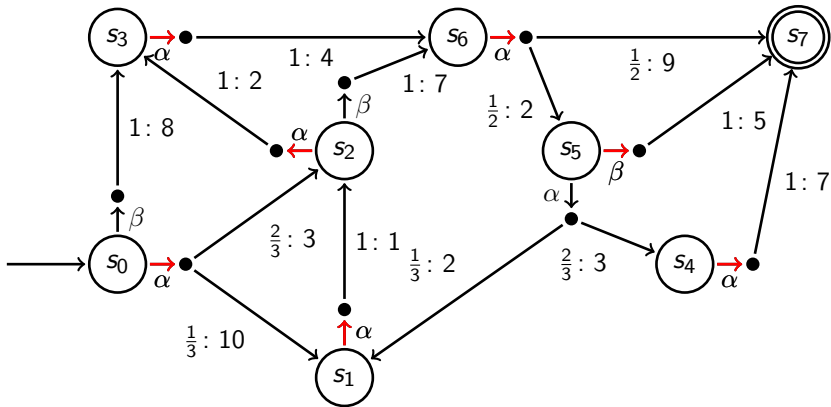


- ▶ Markov Decision Process (MDP)

► Task

- ▶ Give a strategy to reach the goal with minimal *expected* accumulated weights!

- ▶ The *stochastic* shortest path problem



- ▶ Markov Decision Process (MDP)

► Task

- ▶ Give a strategy to reach the goal with minimal *expected* accumulated weights!



## Section 3

### Different variants of the stochastic shortest path problem

## conditional expected accumulated weights

Assume:

- ▶ ...to not reach goal with probability 1

## conditional expected accumulated weights

Assume:

- ▶ ...to not reach goal with probability 1
- ▶ ...having non-negative integer weights

## conditional expected accumulated weights

Assume:

- ▶ ...to not reach goal with probability 1
- ▶ ...having non-negative integer weights
- ▶ ...having two sets of states  $F, G \subseteq S$

## conditional expected accumulated weights

Assume:

- ▶ ...to not reach goal with probability 1
- ▶ ...having non-negative integer weights
- ▶ ...having two sets of states  $F, G \subseteq S$

Maximizing the conditional expected accumulated weight

$$\mathbb{CE} := \mathbb{E}(\Diamond_{\text{goal}} \mid \Diamond_{\text{goal}})$$

## conditional expected accumulated weights

Assume:

- ▶ ...to not reach goal with probability 1
- ▶ ...having non-negative integer weights
- ▶ ...having two sets of states  $F, G \subseteq S$

Maximizing the conditional expected accumulated weight

$$\mathbb{CE} := \mathbb{E}(\Diamond \text{goal} \mid \Diamond \text{goal})$$

$$\mathbb{CE} := \mathbb{E}(\Diamond F \mid \Diamond G)$$

## conditional expected accumulated weights

Assume:

- ▶ ...to not reach goal with probability 1
- ▶ ...having non-negative integer weights
- ▶ ...having two sets of states  $F, G \subseteq S$

Maximizing the conditional expected accumulated weight

$$\mathbb{CE} := \mathbb{E}(\Diamond \text{goal} \mid \Diamond \text{goal})$$

$$\mathbb{CE} := \mathbb{E}(\Diamond F \mid \Diamond G)$$

$$\mathbb{CE}^{max} := \sup_{\mathfrak{G} \in \mathfrak{V}} \mathbb{E}_{\mathcal{M}, s_{init}}^{\mathfrak{G}}(\Diamond F \mid \Diamond G)$$

## conditional expected accumulated weights

Assume:

- ▶ ...to not reach goal with probability 1
- ▶ ...having non-negative integer weights
- ▶ ...having two sets of states  $F, G \subseteq S$

Maximizing the conditional expected accumulated weight

$$\mathbb{CE} := \mathbb{E}(\Diamond \text{goal} \mid \Diamond \text{goal})$$

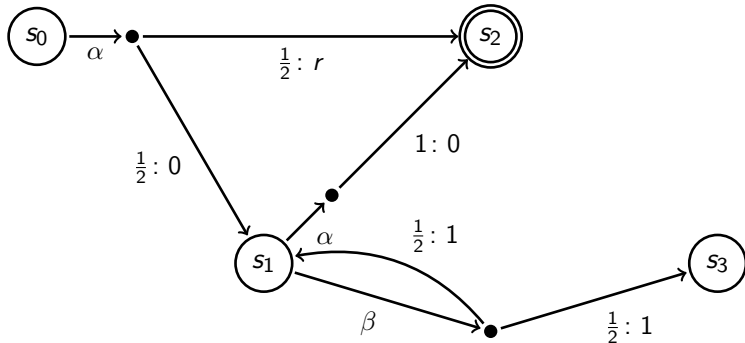
$$\mathbb{CE} := \mathbb{E}(\Diamond F \mid \Diamond G)$$

$$\mathbb{CE}^{max} := \sup_{\mathfrak{G} \in \mathfrak{V}} \mathbb{E}_{\mathcal{M}, s_{init}}^{\mathfrak{G}}(\Diamond F \mid \Diamond G)$$


$$\mathfrak{V} := \{\mathfrak{G} \in \text{Schedulers}(\mathcal{M}) \mid \Pr_{\mathcal{M}, s_{init}}^{\mathfrak{G}}(\Diamond G) > 0 \text{ and} \\ \Pr_{\mathcal{M}, s_{init}}^{\mathfrak{G}}(\Diamond F \mid \Diamond G) = 1\}$$



## conditional expected accumulated reward



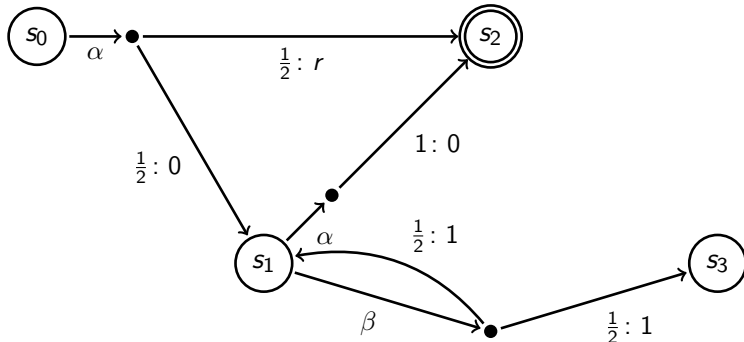
## conditional expected accumulated reward



Finding the best scheduler...

$\mathfrak{S}_n$ : select  $\beta^n \alpha$ ,  $n \in \mathbb{N} \cup \{\infty\}$

## conditional expected accumulated reward



Finding the best scheduler...

$\mathfrak{S}_n$ : select  $\beta^n \alpha$ ,  $n \in \mathbb{N} \cup \{\infty\}$

$$\text{CE} = \frac{\frac{r}{2} + \frac{1}{2} \cdot \frac{1}{2}^n \cdot n}{\frac{1}{2} + \frac{1}{2}^{n+1}}$$

choose scheduler  $\mathfrak{S}_{r+2}$

## complexity: conditional expected accumulated rewards

- ▶ There is a **PTime algorithm** to decide: Is  $\mathbb{CE}^{max}$  finite?

## complexity: conditional expected accumulated rewards

- ▶ There is a **PTime algorithm** to decide: Is  $\mathbb{CE}^{max}$  finite?
- ▶ There is a **pseudo-PTime algorithm** to calculate an upperbound  $\mathbb{CE}^{ub} \geq \mathbb{CE}^{max}$

## complexity: conditional expected accumulated rewards

- ▶ There is a **PTime algorithm** to decide: Is  $\mathbb{CE}^{max}$  finite?
- ▶ There is a **pseudo-PTime algorithm** to calculate an upperbound  $\mathbb{CE}^{ub} \geq \mathbb{CE}^{max}$
- ▶ *If we have  $F = G$  and*  
 $\forall s \in \text{States}(\mathcal{M}) : s \models \exists \Diamond G \Rightarrow \Pr_{\mathcal{M},s}^{min}(\Diamond G) > 0$  *there is a*  
**PTime algorithm** to calculate an upperbound  
 $\mathbb{CE}^{ub} \geq \mathbb{CE}^{max}$

## complexity: conditional expected accumulated rewards

- ▶ There is a **PTime algorithm** to decide: Is  $\mathbb{CE}^{max}$  finite?
- ▶ There is a **pseudo-PTime algorithm** to calculate an upperbound  $\mathbb{CE}^{ub} \geq \mathbb{CE}^{max}$
- ▶ If we have  $F = G$  and  $\forall s \in \text{States}(\mathcal{M}) : s \models \exists \Diamond G \Rightarrow \Pr_{\mathcal{M},s}^{min}(\Diamond G) > 0$  there is a **PTime algorithm** to calculate an upperbound  $\mathbb{CE}^{ub} \geq \mathbb{CE}^{max}$
- ▶ The problem Decide if  $\mathbb{CE}^{max} \bowtie t$  where we have
  - ▶  $t \in \mathbb{Q} \dots$  some rational threshold

## complexity: conditional expected accumulated rewards

- ▶ There is a **PTime algorithm** to decide: Is  $\mathbb{CE}^{max}$  finite?
- ▶ There is a **pseudo-PTime algorithm** to calculate an upperbound  $\mathbb{CE}^{ub} \geq \mathbb{CE}^{max}$
- ▶ If we have  $F = G$  and  $\forall s \in \text{States}(\mathcal{M}) : s \models \exists \Diamond G \Rightarrow \Pr_{\mathcal{M},s}^{min}(\Diamond G) > 0$  there is a **PTime algorithm** to calculate an upperbound  $\mathbb{CE}^{ub} \geq \mathbb{CE}^{max}$
- ▶ The problem Decide if  $\mathbb{CE}^{max} \bowtie t$  where we have
  - ▶  $t \in \mathbb{Q} \dots$  some rational threshold
  - ▶  $\bowtie \in \{<, \leq, \geq, >\}$



# complexity: conditional expected accumulated rewards

- ▶ There is a **PTime algorithm** to decide: Is  $\text{CE}^{max}$  finite?
- ▶ There is a **pseudo-PTime algorithm** to calculate an upperbound  $\text{CE}^{ub} \geq \text{CE}^{max}$
- ▶ If we have  $F = G$  and  $\forall s \in \text{States}(\mathcal{M}) : s \models \exists \Diamond G \Rightarrow \Pr_{\mathcal{M},s}^{min}(\Diamond G) > 0$  there is a **PTime algorithm** to calculate an upperbound  $\text{CE}^{ub} \geq \text{CE}^{max}$
- ▶ The problem Decide if  $\text{CE}^{max} \bowtie t$  where we have
  - ▶  $t \in \mathbb{Q} \dots$  some rational threshold
  - ▶  $\bowtie \in \{<, \leq, \geq, >\}$

is **PSpace-hard**, solvable in **ExpTime** and *for acyclic MDPs* **PSpace-complete**

## complexity: conditional expected accumulated rewards

- ▶ There is a **PTime algorithm** to decide: Is  $\mathbb{CE}^{max}$  finite?
- ▶ There is a **pseudo-PTime algorithm** to calculate an upperbound  $\mathbb{CE}^{ub} \geq \mathbb{CE}^{max}$
- ▶ If we have  $F = G$  and  $\forall s \in \text{States}(\mathcal{M}) : s \models \exists \Diamond G \Rightarrow \Pr_{\mathcal{M},s}^{min}(\Diamond G) > 0$  there is a **PTime algorithm** to calculate an upperbound  $\mathbb{CE}^{ub} \geq \mathbb{CE}^{max}$
- ▶ The problem Decide if  $\mathbb{CE}^{max} \bowtie t$  where we have
  - ▶  $t \in \mathbb{Q} \dots$  some rational threshold
  - ▶  $\bowtie \in \{<, \leq, \geq, >\}$is **PSpace-hard**, solvable in **ExpTime** and *for acyclic MDPs* **PSpace-complete**
- ▶ In **ExpTime** we can compute  $\mathbb{CE}^{max}$  together with an optimal scheduler

# Checking finiteness

- ▶ **PTime** algorithm: given  $\mathcal{M}, F, G$ , two possible outcomes:
  - (1)  $\mathbb{CE}^{max} = \infty$
  - (2)  $\mathbb{CE}^{max}$  is finite, equivalent MDP  $\mathcal{N}$  with two trap states  
*goal, fail*

# Checking finiteness

► **PTime** algorithm: given  $\mathcal{M}, F, G$ , two possible outcomes:

- (1)  $\mathbb{CE}^{max} = \infty$
- (2)  $\mathbb{CE}^{max}$  is finite, equivalent MDP  $\mathcal{N}$  with two trap states  
*goal, fail*

equivalence of  $\mathcal{M}$  and  $\mathcal{N}$ , properties of  $\mathcal{N}$

$$\mathbb{E}_{\mathcal{M}, s_{init}}^{max} (\Diamond F \mid \Diamond G) = \mathbb{E}_{\mathcal{N}, s_{init}}^{max} (\Diamond goal \mid \Diamond goal)$$

# Checking finiteness

► **PTime** algorithm: given  $\mathcal{M}, F, G$ , two possible outcomes:

- (1)  $\mathbb{CE}^{max} = \infty$
- (2)  $\mathbb{CE}^{max}$  is finite, equivalent MDP  $\mathcal{N}$  with two trap states *goal*, *fail*

equivalence of  $\mathcal{M}$  and  $\mathcal{N}$ , properties of  $\mathcal{N}$

$$\mathbb{E}_{\mathcal{M}, s_{init}}^{max} (\Diamond F \mid \Diamond G) = \mathbb{E}_{\mathcal{N}, s_{init}}^{max} (\Diamond goal \mid \Diamond goal)$$

in  $\mathcal{N}$  the state *goal* is reachable from all states  $s \in S_{\mathcal{N}} \setminus \{fail\}$  and

$$\Pr_{\mathcal{N}, s}^{min} (\Diamond (goal \vee fail)) = 1$$

# Checking finiteness

► **PTime** algorithm: given  $\mathcal{M}, F, G$ , two possible outcomes:

- (1)  $\mathbb{CE}^{max} = \infty$
- (2)  $\mathbb{CE}^{max}$  is finite, equivalent MDP  $\mathcal{N}$  with two trap states  
*goal, fail*

equivalence of  $\mathcal{M}$  and  $\mathcal{N}$ , properties of  $\mathcal{N}$

$$\mathbb{E}_{\mathcal{M}, s_{init}}^{max} (\Diamond F \mid \Diamond G) = \mathbb{E}_{\mathcal{N}, s_{init}}^{max} (\Diamond goal \mid \Diamond goal)$$

in  $\mathcal{N}$  the state *goal* is reachable from all states  $s \in S_{\mathcal{N}} \setminus \{fail\}$  and

$$\Pr_{\mathcal{N}, s}^{min} (\Diamond (goal \vee fail)) = 1$$

$\mathcal{N}$  has no *critical* scheduler  $\mathfrak{S}$ :  $\Pr^{\mathfrak{S}}(\Diamond fail) = 1$  and there is a reachable positive  $\mathfrak{S}$ -cycle.

# Threshold Algorithm

We assume that we have such an MDP  $\mathcal{N}$ .

# Threshold Algorithm

We assume that we have such an MDP  $\mathcal{N}$ .

## Observation

- ▶  $\exists$  saturation point  $t \in \mathbb{N}$  such that



# Threshold Algorithm

We assume that we have such an MDP  $\mathcal{N}$ .

## Observation

- ▶  $\exists$  saturation point  $t \in \mathbb{N}$  such that
  - ▶ after  $\pi$  with  $\text{wgt}(\pi) \geq t$  we can rely on a memoryless, deterministic scheduler maximizing the probability to reach *goal*.

# Threshold Algorithm

We assume that we have such an MDP  $\mathcal{N}$ .

## Observation

- ▶  $\exists$  saturation point  $t \in \mathbb{N}$  such that
  - ▶ after  $\pi$  with  $\text{wgt}(\pi) \geq t$  we can rely on a memoryless, deterministic scheduler maximizing the probability to reach *goal*.
  - ▶ Until reaching  $t$  a deterministic reward-based scheduler is sufficient.

# Threshold Algorithm

We assume that we have such an MDP  $\mathcal{N}$ .

## Observation

- ▶  $\exists$  saturation point  $t \in \mathbb{N}$  such that
  - ▶ after  $\pi$  with  $\text{wgt}(\pi) \geq t$  we can rely on a memoryless, deterministic scheduler maximizing the probability to reach *goal*.
  - ▶ Until reaching  $t$  a deterministic reward-based scheduler is sufficient.

## Threshold Algorithm

- ▶ input: MDP  $\mathcal{N}$  as before, threshold  $t \in \mathbb{Q}_{\geq 0}$

# Threshold Algorithm

We assume that we have such an MDP  $\mathcal{N}$ .

## Observation

- ▶  $\exists$  saturation point  $t \in \mathbb{N}$  such that
  - ▶ after  $\pi$  with  $\text{wgt}(\pi) \geq t$  we can rely on a memoryless, deterministic scheduler maximizing the probability to reach *goal*.
  - ▶ Until reaching  $t$  a deterministic reward-based scheduler is sufficient.

## Threshold Algorithm

- ▶ input: MDP  $\mathcal{N}$  as before, threshold  $t \in \mathbb{Q}_{\geq 0}$
- ▶ output:
  - case (1) "no", we do not have  $\mathbb{CE}^{max} > t$
  - case (2) "yes",  $\mathbb{CE}^{max} > t$  and we found a deterministic, reward-based scheduler  $\mathcal{S}$  s.t.

# Threshold Algorithm

We assume that we have such an MDP  $\mathcal{N}$ .

## Observation

- ▶  $\exists$  saturation point  $t \in \mathbb{N}$  such that
  - ▶ after  $\pi$  with  $\text{wgt}(\pi) \geq t$  we can rely on a memoryless, deterministic scheduler maximizing the probability to reach *goal*.
  - ▶ Until reaching  $t$  a deterministic reward-based scheduler is sufficient.

## Threshold Algorithm

- ▶ input: MDP  $\mathcal{N}$  as before, threshold  $t \in \mathbb{Q}_{\geq 0}$
- ▶ output:
  - case (1) "no", we do not have  $\mathbb{CE}^{max} > t$
  - case (2) "yes",  $\mathbb{CE}^{max} > t$  and we found a deterministic, reward-based scheduler  $\mathcal{S}$  s.t.  $\mathcal{S}$  is memoryless after some saturation point.

# What about paths not reaching *goal*?

conditional expectation	partial expectation

## What about paths not reaching *goal*?

conditional expectation	partial expectation
CE	PE

## What about paths not reaching *goal*?

conditional expectation	partial expectation
$\mathbb{CE}$	$\mathbb{PE}$
$\mathbb{CE} = \mathbb{E}(\Diamond \text{goal} \mid \Diamond \text{goal})$	$\pi \not\models \Diamond \text{goal} \Rightarrow \text{wgt}'(\pi) := 0$



## What about paths not reaching *goal*?

$$\oplus goal(\pi) := \begin{cases} \Diamond(\pi), & \pi \models \Diamond F \\ 0, & \text{otherwise, i.e. } \pi \not\models \Diamond F \end{cases}$$

conditional expectation	partial expectation
$\mathbb{C}E$	$\mathbb{P}E$
$\mathbb{C}E = \mathbb{E}(\Diamond goal \mid \Diamond goal)$	$\pi \not\models \Diamond goal \Rightarrow wgt'(\pi) := 0$

## What about paths not reaching *goal*?

$$\oplus goal(\pi) := \begin{cases} \Diamond(\pi), & \pi \models \Diamond F \\ 0, & \text{otherwise, i.e. } \pi \not\models \Diamond F \end{cases}$$

conditional expectation	partial expectation
$\mathbb{CE}$	$\mathbb{PE}$
$\mathbb{CE} = \mathbb{E}(\Diamond goal \mid \Diamond goal)$	$\pi \not\models \Diamond goal \Rightarrow \text{wgt}'(\pi) := 0$
may lead to quite high $\mathbb{CE}$ paired with a low probability of reaching goal	good approximation for maximizing probability of reaching goal and reward until goal

## What about paths not reaching *goal*?

$$\oplus goal(\pi) := \begin{cases} \Diamond(\pi), & \pi \models \Diamond F \\ 0, & \text{otherwise, i.e. } \pi \not\models \Diamond F \end{cases}$$

conditional expectation	partial expectation
$\mathbb{CE}$	$\mathbb{PE}$
$\mathbb{CE} = \mathbb{E}(\Diamond goal \mid \Diamond goal)$	$\pi \not\models \Diamond goal \Rightarrow \text{wgt}'(\pi) := 0$
may lead to quite high $\mathbb{CE}$ paired with a low probability of reaching goal	good approximation for maximizing probability of reaching goal and reward until goal

For both we know...

- **PTime** algorithm to check finiteness of  $\mathbb{CE}^{max}$  ( $\mathbb{PE}^{max}$ )

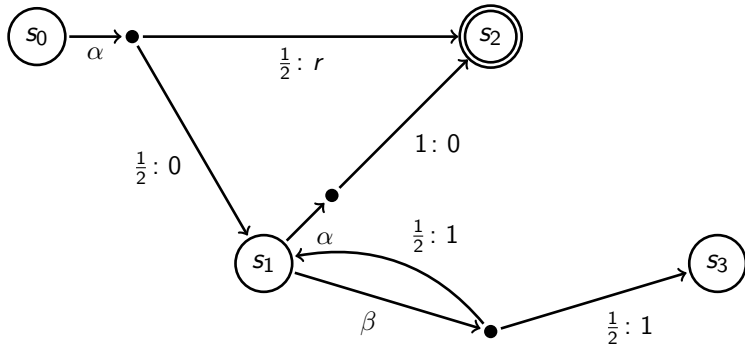
## What about paths not reaching *goal*?

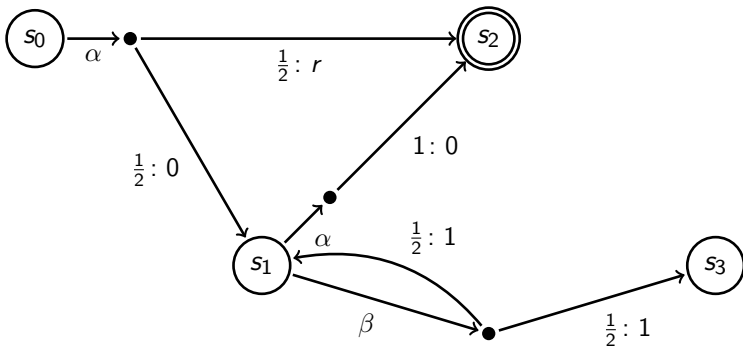
$$\oplus goal(\pi) := \begin{cases} \Diamond(\pi), & \pi \models \Diamond F \\ 0, & \text{otherwise, i.e. } \pi \not\models \Diamond F \end{cases}$$

conditional expectation	partial expectation
$\mathbb{CE}$	$\mathbb{PE}$
$\mathbb{CE} = \mathbb{E}(\Diamond goal \mid \Diamond goal)$	$\pi \not\models \Diamond goal \Rightarrow \text{wgt}'(\pi) := 0$
may lead to quite high $\mathbb{CE}$ paired with a low probability of reaching goal	good approximation for maximizing probability of reaching goal and reward until goal

For both we know...


- ▶ **PTime** algorithm to check finiteness of  $\mathbb{CE}^{max}$  ( $\mathbb{PE}^{max}$ )
- ▶ Both have a saturation point: reward-based schedulers vs. memoryless, det. schedulers





Finding the best scheduler...

$\mathfrak{S}_n$ : select  $\beta^n \alpha$ ,  $n \in \mathbb{N} \cup \{\infty\}$

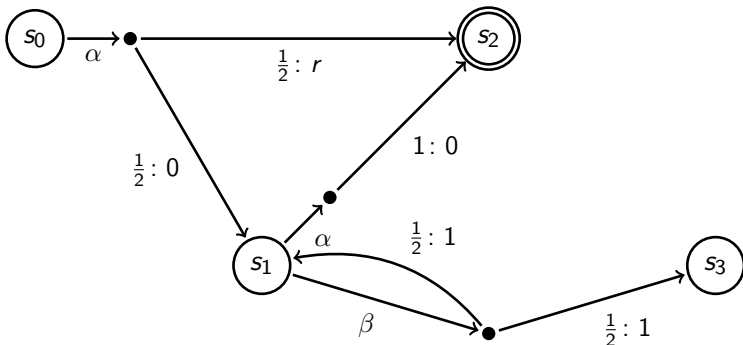


Finding the best scheduler...

$\mathfrak{S}_n$ : select  $\beta^n \alpha$ ,  $n \in \mathbb{N} \cup \{\infty\}$

$$\mathbb{CE} = \frac{\frac{r}{2} + \frac{1}{2} \cdot \frac{1}{2}^n \cdot n}{\frac{1}{2} + \frac{1}{2}^{n+1}}$$

choose scheduler  $\mathfrak{S}_{r+2}$



## Finding the best scheduler...

$\mathfrak{S}_n$ : select  $\beta^n \alpha$ ,  $n \in \mathbb{N} \cup \{\infty\}$

$$\text{CE} = \frac{\frac{r}{2} + \frac{1}{2} \cdot \frac{1}{2}^n \cdot n}{\frac{1}{2} + \frac{1}{2}^{n+1}}$$

choose scheduler  $\mathfrak{S}_{r+2}$

$$\text{PE} = \frac{1}{2}r + \frac{1}{2} \cdot \frac{1}{2}^n \cdot n$$

choose scheduler  $\mathfrak{S}_1$  or  $\mathfrak{S}_2$



## Switching to integer weights...

- ▶ It's not ensured that a saturation point exists

## Switching to integer weights...

- ▶ It's not ensured that a saturation point exists
- ▶ Optimal schedulers may need infinite memory

## Switching to integer weights...

- ▶ It's not ensured that a saturation point exists
- ▶ Optimal schedulers may need infinite memory
- ▶ Optimal values may even become irrational

## Switching to integer weights...

- ▶ It's not ensured that a saturation point exists
- ▶ Optimal schedulers may need infinite memory
- ▶ Optimal values may even become irrational
- ▶ LPs are not anymore sufficient for solving

## Switching to integer weights...

- ▶ It's not ensured that a saturation point exists
- ▶ Optimal schedulers may need infinite memory
- ▶ Optimal values may even become irrational
- ▶ LPs are not anymore sufficient for solving
- ▶ There are approximation methods using LPs

## Section 4

Keep an eye on the variance

# Variance-penalized expectation

Given:

- ▶ MDP  $\mathcal{M}$  with non-negative integer weights

# Variance-penalized expectation

Given:

- ▶ MDP  $\mathcal{M}$  with non-negative integer weights
- ▶ only one trap state *goal* which is reachable from all other states



# Variance-penalized expectation

Given:

- ▶ MDP  $\mathcal{M}$  with non-negative integer weights
- ▶ only one trap state *goal* which is reachable from all other states
- ▶ all states are reachable from  $s_{init}$

# Variance-penalized expectation

Given:

- ▶ MDP  $\mathcal{M}$  with non-negative integer weights
- ▶ only one trap state *goal* which is reachable from all other states
- ▶ all states are reachable from  $s_{init}$

## Variance-penalized expectation

$$\text{VPE}[\lambda]_{\mathcal{M}}^{\mathfrak{S}} := \mathbb{E}_{\mathcal{M}}^{\mathfrak{S}}(\Diamond \text{goal}) - \lambda \cdot \mathbb{V}_{\mathcal{M}}^{\mathfrak{S}}(\Diamond \text{goal})$$

# Variance-penalized expectation

Given:

- ▶ MDP  $\mathcal{M}$  with non-negative integer weights
- ▶ only one trap state *goal* which is reachable from all other states
- ▶ all states are reachable from  $s_{init}$

## Variance-penalized expectation

$$\text{VPE}[\lambda]_{\mathcal{M}}^{\mathfrak{S}} := \mathbb{E}_{\mathcal{M}}^{\mathfrak{S}}(\Diamond \text{goal}) - \lambda \cdot \mathbb{V}_{\mathcal{M}}^{\mathfrak{S}}(\Diamond \text{goal})$$

$$\text{VPE}[\lambda]_{\mathcal{M}}^{\max} := \sup_{\mathfrak{S}} \text{VPE}[\lambda]_{\mathcal{M}}^{\mathfrak{S}}$$

- ▶ Computing a variance-minimal scheduler among all  $\mathbb{E}$ -optimal schedulers is doable in **PTime**. The scheduler can be chosen memoryless.

- ▶ Computing a variance-minimal scheduler among all  $\mathbb{E}$ -optimal schedulers is doable in **PTime**. The scheduler can be chosen memoryless.
- ▶ In **ExpSpace** we can compute  $\mathbb{VPE}^{max}$  and a deterministic scheduler as witness.

- ▶ Computing a variance-minimal scheduler among all  $\mathbb{E}$ -optimal schedulers is doable in **PTime**. The scheduler can be chosen memoryless.
- ▶ In **ExpSpace** we can compute  $\mathbb{VPE}^{max}$  and a deterministic scheduler as witness.
- ▶ The threshold problem, i.e. checking  $\mathbb{VPE}^{max} \geq t$  is decidable in **NExpTime** and is known to be **ExpTime-hard**.

**PTime** algorithm for choosing the best<sub>(variance)</sub> scheduler among all  $\mathbb{E}$ -optimal ones.

1. PTime transformation  $\mathcal{M} \mapsto \mathcal{M}'$  such that

**PTime** algorithm for choosing the best<sub>(variance)</sub> scheduler among all  $\mathbb{E}$ -optimal ones.

1. PTime transformation  $\mathcal{M} \mapsto \mathcal{M}'$  such that
  - ▶  $\mathcal{M}'$  has no 0-ECs



**PTime** algorithm for choosing the best<sub>(variance)</sub> scheduler among all  $\mathbb{E}$ -optimal ones.

1. PTime transformation  $\mathcal{M} \mapsto \mathcal{M}'$  such that
  - ▶  $\mathcal{M}'$  has no 0-ECs
  - ▶ There are mappings from schedulers of  $\mathcal{M}$  to  $\mathcal{M}'$  and vice versa s.t.

**PTime** algorithm for choosing the best<sub>(variance)</sub> scheduler among all  $\mathbb{E}$ -optimal ones.

1. PTime transformation  $\mathcal{M} \mapsto \mathcal{M}'$  such that
  - ▶  $\mathcal{M}'$  has no 0-ECs
  - ▶ There are mappings from schedulers of  $\mathcal{M}$  to  $\mathcal{M}'$  and vice versa s.t.
    - ▶  $\mathbb{E}, \mathbb{V}$  are preserved

**PTime** algorithm for choosing the best<sub>(variance)</sub> scheduler among all  $\mathbb{E}$ -optimal ones.

1. PTime transformation  $\mathcal{M} \mapsto \mathcal{M}'$  such that
  - ▶  $\mathcal{M}'$  has no 0-ECs
  - ▶ There are mappings from schedulers of  $\mathcal{M}$  to  $\mathcal{M}'$  and vice versa s.t.
    - ▶  $\mathbb{E}, \mathbb{V}$  are preserved
2. Transformation  $\mathcal{M}' \mapsto \mathcal{M}''$  such that

**PTime** algorithm for choosing the best<sub>(variance)</sub> scheduler among all  $\mathbb{E}$ -optimal ones.

1. PTime transformation  $\mathcal{M} \mapsto \mathcal{M}'$  such that
  - ▶  $\mathcal{M}'$  has no 0-ECs
  - ▶ There are mappings from schedulers of  $\mathcal{M}$  to  $\mathcal{M}'$  and vice versa s.t.
    - ▶  $\mathbb{E}, \mathbb{V}$  are preserved
2. Transformation  $\mathcal{M}' \mapsto \mathcal{M}''$  such that
  - ▶ All actions not leading to  $\mathbb{E}^{max}$  are removed.

**PTime** algorithm for choosing the best<sub>(variance)</sub> scheduler among all  $\mathbb{E}$ -optimal ones.

1. PTime transformation  $\mathcal{M} \mapsto \mathcal{M}'$  such that
  - ▶  $\mathcal{M}'$  has no 0-ECs
  - ▶ There are mappings from schedulers of  $\mathcal{M}$  to  $\mathcal{M}'$  and vice versa s.t.
    - ▶  $\mathbb{E}, \mathbb{V}$  are preserved
2. Transformation  $\mathcal{M}' \mapsto \mathcal{M}''$  such that
  - ▶ All actions not leading to  $\mathbb{E}^{max}$  are removed.
  - ▶  $\mathcal{M}''$  has no end components

**PTime** algorithm for choosing the best<sub>(variance)</sub> scheduler among all  $\mathbb{E}$ -optimal ones.

1. PTime transformation  $\mathcal{M} \mapsto \mathcal{M}'$  such that
  - ▶  $\mathcal{M}'$  has no 0-ECs
  - ▶ There are mappings from schedulers of  $\mathcal{M}$  to  $\mathcal{M}'$  and vice versa s.t.
    - ▶  $\mathbb{E}, \mathbb{V}$  are preserved
2. Transformation  $\mathcal{M}' \mapsto \mathcal{M}''$  such that
  - ▶ All actions not leading to  $\mathbb{E}^{max}$  are removed.
  - ▶  $\mathcal{M}''$  has no end components
  - ▶ All schedulers of  $\mathcal{M}''$  have equal value for  $\mathbb{E}$ .

## PTime algorithm for choosing the best<sub>(variance)</sub> scheduler among all $\mathbb{E}$ -optimal ones.

1. PTime transformation  $\mathcal{M} \mapsto \mathcal{M}'$  such that
  - ▶  $\mathcal{M}'$  has no 0-ECs
  - ▶ There are mappings from schedulers of  $\mathcal{M}$  to  $\mathcal{M}'$  and vice versa s.t.
    - ▶  $\mathbb{E}, \mathbb{V}$  are preserved
2. Transformation  $\mathcal{M}' \mapsto \mathcal{M}''$  such that
  - ▶ All actions not leading to  $\mathbb{E}^{max}$  are removed.
  - ▶  $\mathcal{M}''$  has no end components
  - ▶ All schedulers of  $\mathcal{M}''$  have equal value for  $\mathbb{E}$ .
3. Solve a system of linear equations to find  $\mathbb{V}^{min}$  + witnessing scheduler

## PTime algorithm for choosing the best<sub>(variance)</sub> scheduler among all $\mathbb{E}$ -optimal ones.

1. PTime transformation  $\mathcal{M} \mapsto \mathcal{M}'$  such that
  - ▶  $\mathcal{M}'$  has no 0-ECs
  - ▶ There are mappings from schedulers of  $\mathcal{M}$  to  $\mathcal{M}'$  and vice versa s.t.
    - ▶  $\mathbb{E}, \mathbb{V}$  are preserved
2. Transformation  $\mathcal{M}' \mapsto \mathcal{M}''$  such that
  - ▶ All actions not leading to  $\mathbb{E}^{max}$  are removed.
  - ▶  $\mathcal{M}''$  has no end components
  - ▶ All schedulers of  $\mathcal{M}''$  have equal value for  $\mathbb{E}$ .
3. Solve a system of linear equations to find  $\mathbb{V}^{min}$  + witnessing scheduler
  - ▶ The scheduler can be chosen memoryless and deterministic



- ▶ Worst-case expected termination times of probabilistic programs

# Applications

- ▶ Worst-case expected termination times of probabilistic programs
- ▶ Finding optimal controls for a motion planning scenario having random external influences

# Applications

- ▶ Worst-case expected termination times of probabilistic programs
- ▶ Finding optimal controls for a motion planning scenario having random external influences
- ▶ Traffic control systems, energy grids

# Applications

- ▶ Worst-case expected termination times of probabilistic programs
- ▶ Finding optimal controls for a motion planning scenario having random external influences
- ▶ Traffic control systems, energy grids
- ▶ Decision making in financial markets

# Applications

- ▶ Worst-case expected termination times of probabilistic programs
- ▶ Finding optimal controls for a motion planning scenario having random external influences
- ▶ Traffic control systems, energy grids
- ▶ Decision making in financial markets

**Limits of applicability?**

## The problem

Given a

- ▶  $\mathbb{Z}$ -weighted MDP  $\mathcal{M}$

# Percentile queries

## The problem

Given a

- ▶  $\mathbb{Z}$ -weighted MDP  $\mathcal{M}$
- ▶ weight threshold  $t \in \mathcal{M}$

## The problem

Given a

- ▶  $\mathbb{Z}$ -weighted MDP  $\mathcal{M}$
- ▶ weight threshold  $t \in \mathcal{M}$
- ▶ probability threshold  $\alpha \in [0, 1] \cap \mathbb{Q}$



# Percentile queries

## The problem

Given a

- ▶  $\mathbb{Z}$ -weighted MDP  $\mathcal{M}$
- ▶ weight threshold  $t \in \mathcal{M}$
- ▶ probability threshold  $\alpha \in [0, 1] \cap \mathbb{Q}$

**Decide:** Is there a scheduler such that  
 $\Pr(\Diamond(\pi) \leq t) \geq \alpha$ ?

## The problem

Given a

- ▶  $\mathbb{Z}$ -weighted MDP  $\mathcal{M}$
- ▶ weight threshold  $t \in \mathcal{M}$
- ▶ probability threshold  $\alpha \in [0, 1] \cap \mathbb{Q}$

**Decide:** Is there a scheduler such that

$$\Pr(\Diamond(\pi) \leq t) \geq \alpha?$$

- ▶ ...PSPACE-hard

## The problem

Given a

- ▶  $\mathbb{Z}$ -weighted MDP  $\mathcal{M}$
- ▶ weight threshold  $t \in \mathcal{M}$
- ▶ probability threshold  $\alpha \in [0, 1] \cap \mathbb{Q}$

**Decide:** Is there a scheduler such that

$\Pr(\Diamond(\pi) \leq t) \geq \alpha$ ?

- ▶ ...**PSpace-hard**
- ▶ ...decidable in **pseudo-PTime**

## The problem

Given a

- ▶  $\mathbb{Z}$ -weighted MDP  $\mathcal{M}$
- ▶ weight threshold  $t \in \mathcal{M}$
- ▶ probability threshold  $\alpha \in [0, 1] \cap \mathbb{Q}$

**Decide:** Is there a scheduler such that

$\Pr(\Diamond(\pi) \leq t) \geq \alpha$ ?

- ▶ ...**PSpace-hard**
- ▶ ...decidable in **pseudo-PTime**
- ▶ ...There is an optimal deterministic scheduler, computable in **ExpTime**

- ▶ Storm model checker
- ▶ Prism model checker